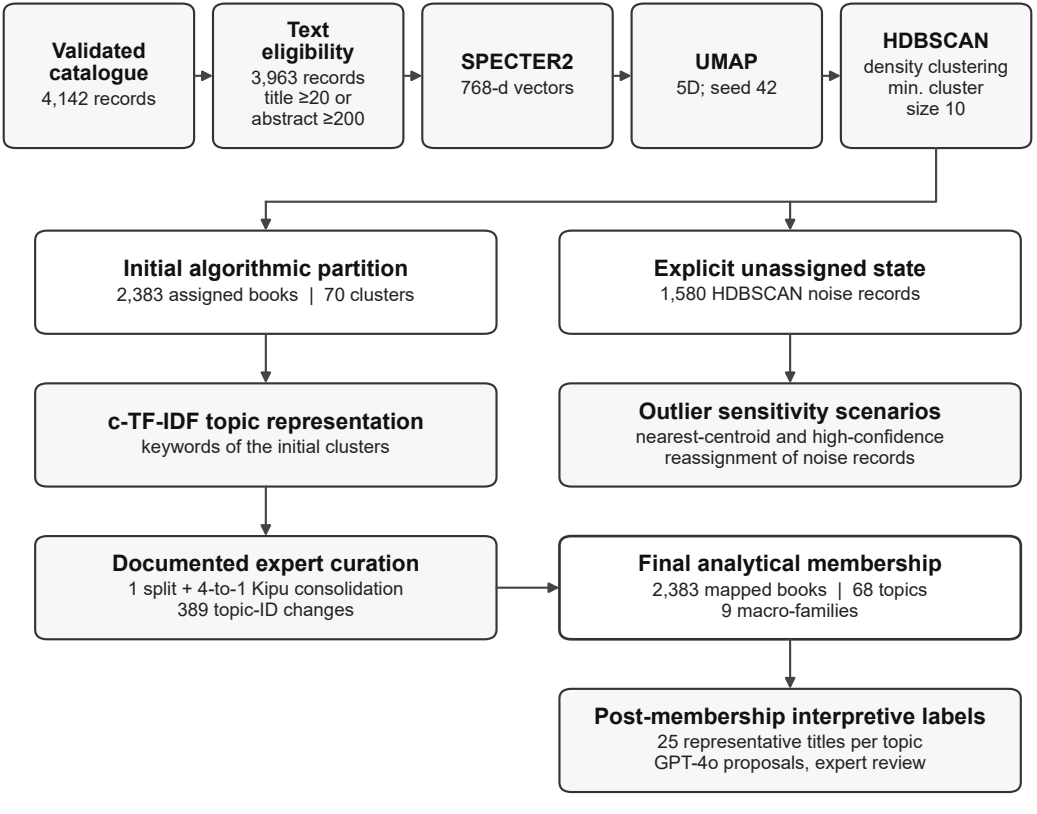
Validated catalogue
4,142 records
Text eligibility
3,963 records
title ≥20 or
abstract ≥200
SPECTER2
768-d vectors
UMAP
5D; seed 42
HDBSCAN
density clustering
min. cluster
size 10
Initial algorithmic partition
2,383 assigned books | 70 clusters
Explicit unassigned state
1,580 HDBSCAN noise records
c-TF-IDF topic representation
keywords of the initial clusters
Outlier sensitivity scenarios
nearest-centroid and high-confidence
reassignment of noise records
Documented expert curation
1 split + 4-to-1 Kipu consolidation
389 topic-ID changes
Final analytical membership
2,383 mapped books | 68 topics
9 macro-families
Post-membership interpretive labels
25 representative titles per topic
GPT-4o proposals, expert review